

1. Linear Embedding - GLoVe

1.1. GLoVe Parameter Count

Since the dimension of \mathbf{w}_i and $\tilde{\mathbf{w}}_i$ are d and the dimension of the biases are 1, there are $2d + 2$ parameters to train for each vocabulary. Thus, we can get that the GLoVe model have $2V(d + 1)$ trainable parameters.

1.2. Expression for gradient $\frac{\partial L}{\partial \mathbf{w}_i}$

According to the given loss function, we can get $\frac{\partial L}{\partial \mathbf{w}_i}$ should be

$$\begin{aligned}\frac{\partial L}{\partial \mathbf{w}_i} &= \frac{\partial}{\partial \mathbf{w}_i} \sum_{i,j=1}^V (\mathbf{w}_i^T \tilde{\mathbf{w}}_j + b_i + \tilde{b}_j - \log X_{ij})^2 \\ &= 2 \sum_{i,j=1}^V (\mathbf{w}_i^T \tilde{\mathbf{w}}_j + b_i + \tilde{b}_j - \log X_{ij}) \tilde{\mathbf{w}}_j\end{aligned}$$

1.3. Implement the gradient update of GLoVe

The code for this question is shown below.

```

1 def grad_GLoVe(W, W_tilde, b, b_tilde, log_co_occurrence):
2     "Return the gradient of GLoVe objective w.r.t W and b."
3     "INPUT: W - Vxd; W_tilde - Vxd; b - Vx1; b_tilde - Vx1; log_co_occurrence: VxV"
4     "OUTPUT: grad_W - Vxd; grad_W_tilde - Vxd; grad_b - Vx1; grad_b_tilde - Vx1"
5     n, _ = log_co_occurrence.shape
6
7     if not W_tilde is None and not b_tilde is None:
8         ##### YOUR CODE HERE #####
9         loss = (W @ W_tilde.T + b @ np.ones([1, n]) + np.ones([n, 1]) @ b_tilde.T -
10              log_co_occurrence)
11         grad_W = 2 * (loss @ W_tilde)
12         grad_W_tilde = 2 * (loss.T @ W)
13         grad_b = 2 * (np.ones([1, n]) @ loss).T
14         grad_b_tilde = 2 * (np.ones([n, 1]) @ loss).T
15     else:
16         loss = (W @ W.T + b @ np.ones([1, n]) + np.ones([n, 1]) @ b.T - 0.5 * (
17              log_co_occurrence + log_co_occurrence.T))
18         grad_W = 4 * (W.T @ loss).T
19         grad_W_tilde = None
20         grad_b = 4 * (np.ones([1, n]) @ loss).T
21         grad_b_tilde = None
22
23     return grad_W, grad_W_tilde, grad_b, grad_b_tilde

```

1.4. Effects of embedding dimension

2. Network architecture

2.1. Number of parameters in neural network model

Let's find the number of the trainable parameters of the word embedding weights first. Since there are V words in the dictionary and the dimension of the word embedding layer is $N \times D$, we can get that there are $V \times D$ trainable parameters.

For weights between the word embedding layer and the hidden layer, since there are H units in the hidden layer, the dimension of the matrix that connects two layers should be $ND \times H$, which is the number of trainable parameters.

Thus, for the biases of the hidden layer, there should be $H \times 1$ trainable parameters.

Similarly, since the output layer consists of V words, there are $V \times H$ trainable parameters for

weights between the hidden layer and the output layer.

Thus, for the biases of the output layer, there should be $V \times 1$ trainable parameters.

Since V is much larger than other variables, we only need to consider the part that depends on V .

Thus, we can get that weights between the hidden layer and the output layer, *hid_to_output_weights*, has the largest number of trainable parameters since $H > D$.

2.2. Number of parameters in n-gram mode

For each gram, we can choose any word from V words. Thus, there are V^N number of combinations for the previous N words. For the prediction, since the output layer is a softmax over the V words, there are V words. Thus, there are V^{N+1} entries in the n -gram model scale with N .

2.3. Comparing neural network and n-gram model scaling

3. Training the Neural Network

3.1. Implement gradient with respect to output layer inputs

The code for this question is shown below.

```

1 def compute_loss_derivative(self, output_activations, expanded_target_batch,
  target_mask):
2     """Compute the derivative of the multiple target position cross-entropy
  loss function \n"
3
4     For example:
5
6     [y_{0} ... y_{V-1}] [y_{V}, ..., y_{2*V-1}] [y_{2*V} ... y_{i,3*V-1}] [
  y_{3*V} ... y_{i,4*V-1}]
7
8     Where for column j + n*V,
9
10    y_{j + n*V} = e^{z_{j + n*V}} / \sum_{m=0}^{V-1} e^{z_{m + n*V}}, for
  n=0,...,N-1
11
12    This function should return a dC / dz matrix of size [batch_size x (
  vocab_size * context_len)],
13    where each row i in dC / dz has columns 0 to V-1 containing the gradient
  the 1st output
14    context word from i-th training example, then columns vocab_size to 2*
  vocab_size - 1 for the 2nd
15    output context word of the i-th training example, etc.
16
17    C is the loss function summed acrossed all examples as well:
18
19    C = -\sum_{i,j,n} mask_{i,n} (t_{i, j + n*V} log y_{i, j + n*V}), for
  j=0,...,V, and n=0,...,N
20
21    where mask_{i,n} = 1 if the i-th training example has n-th context word as
  the target,
22    otherwise mask_{i,n} = 0.
23
24    The arguments are as follows:
25
26    output_activations - A [batch_size x (context_len * vocab_size)]
  tensor,
27    for the activations of the output layer, i.e. the y_j's.
28    expanded_target_batch - A [batch_size (context_len * vocab_size)]
  tensor,
29    where expanded_target_batch[i,n*V:(n+1)*V] is the indicator vector
  for
30    the n-th context target word position, i.e. the (i, j + n*V) entry
  is 1 if the
31    i'th example, the context word at position n is j, and 0 otherwise
  .
32    target_mask - A [batch_size x context_len x 1] tensor, where
  target_mask[i,n] = 1
33    if for the i'th example the n-th context word is a target position
  , otherwise 0

```

```

34
35     Outputs:
36         loss_derivative — A [batch_size x (context_len * vocab_size)] matrix,
37         where loss_derivative[i,0:vocab_size] contains the gradient
38         dC / dz_0 for the i-th training example gradient for 1st output
39         context word, and loss_derivative[i,vocab_size:2*vocab_size] for
40         the 2nd output context word of the i-th training example, etc.
41     """
42
43     ##### YOUR CODE HERE
44     #####
45     # Loss
46     loss = output_activations - expanded_target_batch
47     expanded_mask = np.repeat(target_mask, self.vocab_size, axis=1).reshape(
48     loss.shape)
49     return np.multiply(expanded_mask, loss)
50     #
51     #####

```

3.2. Implement gradient with respect to parameters

The code for this question is shown below.

```

1 def back_propagate(self, input_batch, activations, loss_derivative):
2     """Compute the gradient of the loss function with respect to the trainable
3     parameters
4     of the model. The arguments are as follows:
5
6         input_batch — the indices of the context words
7         activations — an Activations class representing the output of Model.
8         compute_activations
9         loss_derivative — the matrix of derivatives computed by
10        compute_loss_derivative
11
12        Part of this function is already completed, but you need to fill in the
13        derivative
14        computations for hid_to_output_weights_grad, output_bias_grad,
15        embed_to_hid_weights_grad,
16        and hid_bias_grad. See the documentation for the Params class for a
17        description of what
18        these matrices represent."""
19
20        # The matrix with values dC / dz_j, where dz_j is the input to the jth
21        hidden unit,
22        # i.e. h_j = 1 / (1 + e^{-z_j})
23        hid_deriv = np.dot(loss_derivative, self.params.hid_to_output_weights) \
24        * activations.hidden_layer * (1. - activations.hidden_layer)
25
26        ##### YOUR CODE HERE
27        #####
28        hid_to_output_weights_grad = loss_derivative.T @ activations.hidden_layer
29        output_bias_grad = np.sum(loss_derivative, axis=0)
30        embed_to_hid_weights_grad = hid_deriv.T @ activations.embedding_layer
31        hid_bias_grad = np.sum(hid_deriv, axis=0)
32        #
33        #####
34
35        # The matrix of derivatives for the embedding layer
36        embed_deriv = np.dot(hid_deriv, self.params.embed_to_hid_weights)
37
38        # Embedding layer
39        word_embedding_weights_grad = np.zeros((self.vocab_size, self.
40        embedding_dim))
41        for w in range(self.context_len):
42            word_embedding_weights_grad += np.dot(self.indicator_matrix(
43            input_batch[:, w:w+1], mask_zero_index=False).T,
44            embed_deriv[:, w * self.
45            embedding_dim:(w + 1) * self.embedding_dim])

```

```

34
35         return Params(word_embedding_weights_grad, embed_to_hid_weights_grad,
36                        hid_to_output_weights_grad,
37                        hid_bias_grad, output_bias_grad)

```

3.3. Print the gradients The output for `print_gradients()` is shown below.

```

1 loss_derivative[2, 5] 0.0
2 loss_derivative[2, 121] 0.0
3 loss_derivative[5, 33] 0.0
4 loss_derivative[5, 31] 0.0
5
6 param_gradient.word_embedding_weights[27, 2] 0.0
7 param_gradient.word_embedding_weights[43, 3] 0.011596892511489458
8 param_gradient.word_embedding_weights[22, 4] -0.0222670623817297
9 param_gradient.word_embedding_weights[2, 5] 0.0
10
11 param_gradient.embed_to_hid_weights[10, 2] 0.3793257091930164
12 param_gradient.embed_to_hid_weights[15, 3] 0.01604516132110917
13 param_gradient.embed_to_hid_weights[30, 9] -0.4312854367997419
14 param_gradient.embed_to_hid_weights[35, 21] 0.06679896665436337
15
16 param_gradient.hid_bias[10] 0.023428803123345148
17 param_gradient.hid_bias[20] -0.024370452378874197
18
19 param_gradient.output_bias[0] 0.000970106146902794
20 param_gradient.output_bias[1] 0.16868946274763222
21 param_gradient.output_bias[2] 0.0051664774143909235
22 param_gradient.output_bias[3] 0.15096226471814364
23

```

3.4. Run model training

4. Arithmetics and Analysis

4.1. t-SNE




Figure 1: The tsne plot representation using the trained weights.




Figure 2: The tsne plot GLoVE representation.




Figure 3: The 2d GLoVe representation.




Figure 4: The tsne plot GLoVe representation.

According to Figure 1 (tsne_plot_representation), we can see that words with the similar function in sentences or the structure in terms of the speech gather together. For example, on the top left of the plot, there is a cluster of auxiliary verbs including “would”, “could”, and “should”, and a cluster of question words such as “what”, “how”, and “who” on the middle left of the plot.

For Figure 1 and 2 (tsne_plot_GLoVe_representation), we can see that words in Figure 1 are distributed like a linear on the main diagonal, while words in Figure 2 diverge circularly around the biggest cluster on the lower middle. In addition, we found that positions of different clusters are different in two plots.

Comparing with Figure 1 and 2, we can see that words in Figure 3 (plot_2d_GLoVe_representation) seem to be more clustered and are distributed like a fan shape. Many nouns gather together on the bottom left cluster, which is the largest cluster and words are fan out upward.

4.2. Word Analogy Arithmetic

4.2.1. Specific example

The results are shown below.

```

1 ## GLoVe embeddings
2 The top 10 closest words to emb(he) - emb(him) + emb(her) are:
3 he: 1.4213098857979793
4 she: 1.48167433432594
5 said: 2.1025960106397767
6 then: 2.2720425987761406
7 does: 2.301964867719902
8 says: 2.318047293286045
9 who: 2.328984314854128
10 where: 2.334702431567161
11 did: 2.353623598835888

```

```

12 should: 2.4126428205989865
13
14 # Concatenation of W_final_asym, W_tilde_final_asym
15 The top 10 closest words to emb(he) - emb(him) + emb(her) are:
16 he: 2.046826000951795
17 she: 2.3455038844018743
18 i: 3.0624522787351487
19 we: 3.2848647174761094
20 they: 3.390910580609287
21 you: 4.568945007203308
22 john: 4.805241000654006
23 program: 5.084420284826234
24 president: 5.104152796566877
25 never: 5.111163924178705
26
27 # Averaging asymmetric GLoVE vectors
28 The top 10 closest words to emb(he) - emb(him) + emb(her) are:
29 he: 1.0154702232698416
30 she: 1.0744126028585648
31 should: 1.6078139338035942
32 could: 1.6799073061855805
33 i: 1.693953840244398
34 would: 1.700020962168106
35 did: 1.766206937557185
36 can: 1.7744377144463797
37 might: 1.7765376616413824
38 will: 1.7931360498829227
39
40 ### Neural Network Word Embeddings
41 The top 10 closest words to emb(he) - emb(him) + emb(her) are:
42 he: 2.4284684644619032
43 she: 17.4415802699889
44 have: 25.921497697983263
45 they: 25.981587972296392
46 want: 26.437644546989542
47 we: 27.128094534488834
48 i: 27.215833550319473
49 but: 28.03028938337095
50 about: 28.163403568035555
51 this: 28.531350495330678
52

```

According to the outputs, we can see that the closest word that is not “he”, “him”, or “her” is “she” for all 4 different arithmetic. The corresponding distances are shown above. According to four plots, they all show the parallelogram property of the quadruplets approximately. Figure 5, 6, 7, and 8 show how they present in the corresponding plot.




Figure 5: The tsne plot representation using the trained weights.




Figure 6: The tsne plot GloVe representation.




Figure 7: The 2d GloVe representation.




Figure 8: The tsne plot GloVe representation.